


# XGBoost

Efficient boosting with tree models

Martin Jullum

Big Insight lunch, Jan 31, 2018



# XGBoost = eXtreme Gradient Boosting


- ▶ A machine learning library built around an efficient implementation of boosting for tree models (like GBM)
  - Developed by Tianqi Chen (Uni. Washington) i 2014
- ▶ Core library in C++, with interfaces for many languages/platforms
  - C++, Python, R, Julia, Java, etc.
  - Distributed version for Hadoop + Spark
- ▶ Engineering goal: “Push the limit of computational resources for boosted tree algorithms”
  - Parallelizable, cheap on memory, scales to large data sets
- ▶ Very powerful and flexible – lots of (hyper)parameters
- ▶ Huge success
  - «Winning practically every prediction competition on Kaggle»

# Problem setup

- ▶ Assume we have training data set of size  $n$ 
  - Response:  $y_i$
  - Covariates:  $x_i = (x_{i1}, \dots, x_{ip})^\top, i = 1, \dots, n$
- ▶ Want to train a model  $f$  on these data such that  $f(x_i)$  approximates  $y_i$  as well as possible (on a separate test data set!) in terms of a loss function  $L(y, f)$


# Tree models (I)

- Conceptually possibly the simplest statistical model existing!
  - The function is evaluated by a series of conditional IF-ELSE rules
  - As a tree: Start at the root and work your way through the branches depending on your covariate values, ending up at the leaves



Trained tree model for survival rate on Titanic

# Tree models (II)



3 visualizations of the same tree model

- May be written as a weighted sum of indicator values

$$f(x) = \sum_{j=1}^T \theta_j 1_{\{x \in R_j\}}$$

# Training a tree model

► Computationally intractable to find the best partitioning w.r.t. general  $L(y, f)$ , so need a greedy algorithm, which iteratively grows the tree

► Algorithm:

- For each leaf node  $N_j, j = 1, \dots$ , in the current tree DO:
  - For each covariate  $x_j$ , find the split point corresponding to new potential regions  $R_{1j}, R_{2j}$  minimizing the split loss

$$\sum_{i \in N_j} [L(y_i, \hat{y}_{R_{1j}}) + L(y_i, \hat{y}_{R_{2j}})]$$

where  $\hat{y}_{R_{kj}} = \operatorname{argmin}_c \sum_{i \in R_{kj}} L(y_i, c)$ .

- Choose the leaf node, covariate and split point with smallest split loss
- Perform the split if loss reduction is large enough in terms of e.g. previous loss reductions, depth, number of nodes, etc.
- REPEAT

# Properties of tree models

## ► Benefits

- Models non-linearities and interactions directly
- Invariant under monotone transformations of the covariates
- Easy to train – scales well to large data sets
- Naturally combines continuous and categorical data
- Easy to explain and interpret
- Can handle missing data
- Robust to outliers in the covariates

## ► Drawbacks

- Limited predictive power
- High variance
- Somewhat arbitrary handling of overfitting/regularization
- Lack smoothness


# Boosting: The principle

- ▶ Kearns (1988) asked whether weak classifiers could be combined into a strong classifier
- ▶ Freund and Schapire (1997): YES, with AdaBoost
- ▶ AdaBoost idea
  - Iteratively fit simple models  $f_m(x)$  (weak learners) trying to correct «mistakes» of previous models
  - Combine them additively into a model ensemble with good predictive performance (strong learner)

$$f^{(M)}(x) = \sum_{m=1}^M f_m(x)$$



# Example Adaboost



# Adaboost as FSAM

- ▶ Friedman et al. (2000): Adaboost is equivalent to Forward Stagewise Additive Modelling (FSAM) with the loss function:  
 $L(y, f(x)) = \exp(-yf(x))$
- ▶ FSAM: For  $m = 1, \dots, M$ , find model  $f_m$  by minimizing the empirical risk

$$f_m = \underset{h \in \Phi}{\operatorname{argmin}} \frac{1}{n} \sum_{i=1}^n L(y_i, f^{(m-1)}(x_i) + h(x_i))$$

- $f^{(m-1)}(x) = \sum_{j=1}^{m-1} f_j(x)$ , with  $f^{(0)}(x) = 0$
- For some model class  $\Phi$
- A very general procedure, but hard to do for general loss function

# Gradient boosting


## ► Gradient descent

- Iterative procedure for finding minimum of (multivariate) function  $s(z)$
- Iteratively take steps along the negative gradient:  $z_m = z_{m-1} - \rho_m s'(z_{m-1})$

## ► Notation: Let $s_i(z) = L(y_i, z)$

## ► Gradient boosting (Friedman (2001))

- Take a gradient descent step towards the minimum of  $s_i$ , at  $z = f^{(m-1)}(x_i)$  for all  $i$
- Class restriction solved by using the function closest in L2 to the negative gradient:  $f_{m,0} = \underset{h \in \Phi}{\operatorname{argmin}} \sum_{i=1}^n [-s'_i(f^{(m-1)}(x_i)) - h(x_i)]^2$
- The step length is found by 
$$\rho_m = \underset{\rho}{\operatorname{argmin}} \sum_{i=1}^n L(y_i, f^{(m-1)}(x_i) + \rho f_{m,0}(x_i)),$$
- Finally:  $f_m(x) = \gamma \rho_m f_{m,0}(x)$ , for some pre-set learning rate  $\gamma \in (0,1]$
- This is the most common boosting method, e.g. gbm package in R



## 2. order approximation

- ▶ Approximate  $L(y_i, f^{(m-1)}(x_i) + h(x_i))$ , using a 2. order Taylor approximation of  $s_i(z)$  around  $z = f^{(m-1)}(x_i)$
- ▶ 
$$s^*(f^{(m-1)}(x_i) + h(x_i)) = s_i(f^{(m-1)}(x_i)) + s_i'(f^{(m-1)}(x_i))h(x_i) + \frac{1}{2}s_i''(f^{(m-1)}(x_i))h(x_i)^2$$
- ▶ Inserting this into the FSAM solution gives
- ▶ 
$$\begin{aligned} f_{m,0} &= \underset{h \in \Phi}{\operatorname{argmin}} \sum_{i=1}^n [s^*(f^{(m-1)}(x_i) + h(x_i))] \\ &= \underset{h \in \Phi}{\operatorname{argmin}} \sum_{i=1}^n [s_i'(f^{(m-1)}(x_i))h(x_i) + \frac{1}{2}s_i''(f^{(m-1)}(x_i))h(x_i)^2] \\ &= \underset{h \in \Phi}{\operatorname{argmin}} \sum_{i=1}^n \frac{1}{2} s_i''(f^{(m-1)}(x_i)) \left[ -\frac{s_i'(f^{(m-1)}(x_i))}{s_i''(f^{(m-1)}(x_i))} - h(x_i) \right]^2 \end{aligned}$$
- ▶ Just a weighted least squares problem w.r.t.  $h \in \Phi$
- ▶ Finally:  $f_m(x) = \gamma f_{m,0}(x)$ , for some pre-set learning rate  $\gamma$
- ▶ This is the method used by XGBoost, with  $\Phi$  being tree models (originally proposed through LogitBoost)

# XGBoost – methodological improvements

- ▶ Tree boosting inherits most benefits and fixes the drawbacks of individual tree models
- ▶ 2. order approx. to FSAM – more precise than regular gradient boosting
- ▶ Introduced regularization directly in the tree growing procedure
  - Actually tries to minimize,  $L(y_i, f^{(m-1)}(x_i) + h(x_i)) + \Omega(h)$ ,
  - $\Omega(h) = \gamma T + \frac{1}{2} \lambda \sum_j^T w_j^2 + \alpha \sum_j^T |w_j|$ , for  $w_j$  the leaf values of tree of depth  $T$
  - Also other regularization parameters available
- ▶ Subsampling of both rows and columns of covariate matrix available for better generalization properties

# XGBoost – technical improvements

- ▶ Very fast and cheap on memory
  - Store data in internal sparsity aware format – memory friendly
  - The tree learning algorithm utilizes the sparse structure
  - Parallelizes tree learning per covariate
  - Example:  $n=2 \times 10^6$ ,  $p=200$ ,  $Y=\{0,1\}$ ,  $\text{depth}=6$ , **150sec** with 16 threads, a few GB of RAM consumption.
- ▶ Allows the user to view the performance of the current model during training
- ▶ Can automatically stop boosting when performance on separate validation set decreases
- ▶ User can set custom loss function and evaluation metric for stopping
- ▶ Implemented direct handling of missing values – learning a default direction for NA

# Some recent community contributions

- ▶ DART (Dropout Additive Regression Trees) (Feb 2016)
  - Drop given proportion of trained trees when learning a new tree
  - More randomization -> link to random forest
- ▶ Histogram approach (Jan 2017)
  - Discretize continuous covariates into default bins for faster training, 4-10 times faster
- ▶ GPU version (Aug 2017)
  - 2-4 times faster than histogram approach on CPU
- ▶ Covariate contribution per prediction supported natively by SHAP (Oct 2017)

# Remarks

- ▶ Competitors
  - LightGBM (Microsoft)
    - Very similar, not as mature and feature rich
    - Slightly faster than XGBoost – much faster when it was published
  - CatBoost (Yandex, “Russian Google”)
    - Also similar, but handles categorical variables directly
    - Benchmarks show better results, but is much slower
- ▶ CRAN verison of xgboost is outdated (do NOT use!)
  - Install from private repo (simple) or directly from Github (advanced)  
<http://xgboost.readthedocs.io/en/latest/build.html#r-package-installation>
- ▶ Can be called from caret, h2o R-packages + scikit-learn in Python
- ▶ I have still not seen an example where Random Forest outperforms XGBoost



# Key resources

- ▶ Didrik Nielsen, Master thesis NTNU, 2016:  
<https://brage.bibsys.no/xmlui/handle/11250/2433761>
- ▶ Chen & Guestrin (2016), XGBoost: A Scalable Tree Boosting System: <https://arxiv.org/abs/1603.02754>
- ▶ Hastie et al. (2009), Elements of Statistical Learning, Ch 9.2 + 10
- ▶ XGBoost Github: <https://github.com/dmlc/xgboost>
- ▶ XGBoost documentation: <http://xgboost.readthedocs.io>
- ▶ Slides from Meetup in LA with Tianqi Chen:  
<http://datascience.la/xgboost-workshop-and-meetup-talk-with-tianqi-chen/>